

CISC 6930 Course Project

December 7, 2018

1 UCI Diabetes - Readmission in 3 months or not?

In this project, several models will be created from the UCI diabetes dataset that will predict whether a patient is likely to be readmitted to the hospital for diabetes in 30 days or not.

1.1 Import packages

```
In [1]: # Import packages
import pandas as pd
import numpy as np
import scipy.stats as ss
import seaborn as sns
import matplotlib.pyplot as plt
import time

from sklearn.preprocessing import Normalizer

from sklearn.feature_selection import chi2, SelectPercentile
from imblearn.over_sampling import SMOTE

from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold, cross_val_predict

from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier, ExtraTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import LinearSVC
from sklearn.linear_model import Perceptron

from sklearn.ensemble import RandomForestClassifier, VotingClassifier

from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix, classification_report

import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
```

2 Data Analysis

Start by reading in the data and removing extraneous information of individual patients.

```
In [2]: # Read in dataset
df = pd.read_csv("diabetic_data.csv", na_values = "?", low_memory = False)

# Drop observation specific features
df = df.drop(["encounter_id", "patient_nbr"], axis = 1)

# Display class label counts
df.readmitted.value_counts()
```

```
Out[2]: NO      54864
      >30      35545
      <30      11357
      Name: readmitted, dtype: int64
```

```
In [3]: # Adjust to binary classification problem
df.readmitted = df.readmitted.apply(lambda x: 1 if x == '<30' else 0)
```

```
In [4]: # Check class labels
print("After adjusting the class labels, \n",
      'total number of class 0 (NO and >30) instances:',
      len(df.loc[df['readmitted'] == 0, 'readmitted']), "\n",
      "total number of class 1 (< 30) instances:",
      len(df.loc[df['readmitted'] == 1, 'readmitted']))
```

```
After adjusting the class labels,
total number of class 0 (NO and >30) instances: 90409
total number of class 1 (< 30) instances: 11357
```

Note which columns have continuous data and which have categorical data. If a certain column is of a wrong type, change it to the other type.

```
In [5]: # List the type of feature each column is
df.columns.groupby(df.dtypes)
```

```
Out[5]: {dtype('int64'): Index(['admission_type_id', 'discharge_disposition_id', 'admission_source_id',
                                'time_in_hospital', 'num_lab_procedures', 'num_procedures',
                                'num_medications', 'number_outpatient', 'number_emergency',
                                'number_inpatient', 'number_diagnoses', 'readmitted'],
                                dtype='object'),
         dtype('O'): Index(['race', 'gender', 'age', 'weight', 'payer_code', 'medical_specialty',
                             'diag_1', 'diag_2', 'diag_3', 'max_glu_serum', 'A1Cresult', 'metformin',
                             'repaglinide', 'nateglinide', 'chlorpropamide', 'glimepiride',
                             'acetohexamide', 'glipizide', 'glyburide', 'tolbutamide',
                             'pioglitazone', 'rosiglitazone', 'acarbose', 'miglitol', 'troglitazone',
```

```
'tolazamide', 'examide', 'citoglipton', 'insulin',
'glyburide-metformin', 'glipizide-metformin',
'glimepiride-pioglitazone', 'metformin-rosiglitazone',
'metformin-pioglitazone', 'change', 'diabetesMed'],
dtype='object'))}
```

IDs are to be treated as categorical data rather than as numbers. The 3 ID columns are mistaken for numerical columns. Therefore change those to string types.

```
In [6]: # Change ID columns from numerical to string type
df.admission_type_id = df.admission_type_id.astype('str')
df.discharge_disposition_id = df.discharge_disposition_id.astype('str')
df.admission_source_id = df.admission_source_id.astype('str')
```

2.1 Exploratory Data Analysis

```
In [7]: # What is the size of the dataset
df.shape
```

```
Out[7]: (101766, 48)
```

There are 101,766 records in the dataset. Explore the dataset. If a certain column is overly populated by a frequent value, over 98%, remove the column from the dataset.

To find heavy tailed distributions, first find a cutoff point.

```
In [8]: # Find cut off point
0.98 * df.shape[0]
```

```
Out[8]: 99730.68
```

Find the columns where at least 99,731 rows have the same value.

```
In [9]: # Get list of features ordered by top most frequent values
df.describe(exclude = ['int64']).T.sort_values(by = ['freq'], ascending = False)
```

```
Out[9]:
```

	count	unique	top	freq
examide	101766	1	No	101766
citoglipton	101766	1	No	101766
glimepiride-pioglitazone	101766	2	No	101765
acetoexamide	101766	2	No	101765
metformin-pioglitazone	101766	2	No	101765
metformin-rosiglitazone	101766	2	No	101764
troglitazone	101766	2	No	101763
glipizide-metformin	101766	2	No	101753
tolbutamide	101766	2	No	101743
miglitol	101766	4	No	101728
tolazamide	101766	3	No	101727
chlorpropamide	101766	4	No	101680
acarbose	101766	4	No	101458

nateglinide	101766	4	No	101063
glyburide-metformin	101766	4	No	101060
repaglinide	101766	4	No	100227
glimepiride	101766	4	No	96575
max_glu_serum	101766	4	None	96420
rosiglitazone	101766	4	No	95401
pioglitazone	101766	4	No	94438
glyburide	101766	4	No	91116
glipizide	101766	4	No	89080
A1Cresult	101766	4	None	84748
metformin	101766	4	No	81778
diabetesMed	101766	2	Yes	78363
race	99493	5	Caucasian	76099
discharge_disposition_id	101766	26	1	60234
admission_source_id	101766	17	7	57494
change	101766	2	No	54755
gender	101766	3	Female	54708
admission_type_id	101766	8	1	53990
insulin	101766	4	No	47383
payer_code	61510	17	MC	32439
age	101766	10	[70-80)	26068
medical_specialty	51817	72	InternalMedicine	14635
diag_3	100343	789	250	11555
diag_1	101745	716	428	6862
diag_2	101408	748	276	6752
weight	3197	9	[75-100)	1336

The 16 columns that are over 98% saturated with one value are: *examide*, *citoglipton*, *glimepiride-pioglitazone*, *acetohexamide*, *metformin-pioglitazone*, *metformin-rosiglitazone*, *trogliptazone*, *glipizide-metformin*, *tolbutamide*, *miglitol*, *tolazamide*, *chlorpropamide*, *acarbose*, *nateglinide*, *glyburide-metformin* and *repaglinide*. Remove those from the dataframe.

```
In [10]: # Drop the overly saturated features
df = df.drop(['examide', 'citoglipton', 'glimepiride-pioglitazone',
              'acetohexamide', 'metformin-pioglitazone',
              'metformin-rosiglitazone', 'trogliptazone',
              'glipizide-metformin',
              'tolbutamide', 'miglitol', 'tolazamide',
              'chlorpropamide',
              'acarbose', 'nateglinide', 'glyburide-metformin',
              'repaglinide'], axis = 1)
```

```
In [11]: # What is the new length of the dataset?
df.shape
```

```
Out[11]: (101766, 32)
```

An important note about this dataset is that there are 3 columns designated for diagnoses. These columns give the icd9 code for what the patient was diagnosed with on the primary and

two secondary diagnoses. According to the research article that was published alongside with the dataset, these diagnoses can be coded into a larger group for our use. These groups are circulatory, respiratory, digestive, diabetes, injury, musculoskeletal, genitourinary, neoplasms, symptoms, skin/tissue, parasitic and others. The others group makes up 17.3% of the dataset and is broken down into 7 smaller groups each of having 2.2% or less occurrence in the dataset.

Source: <https://www.hindawi.com/journals/bmri/2014/781670/>

Thus code the diagnoses columns into these groups.

```
In [12]: # A function to categorize diagnosis icd9 codes
def diag_group(df, col):
    new_diags = []
    for i in df[col]:
        if pd.isnull(i): new_diags.append(np.nan)
        else:
            try:
                val = float(i)
                if (val >= 390 and val <= 460) or val == 785:
                    new_diags.append("circulatory")
                elif (val >= 460 and val <= 520) or val == 786:
                    new_diags.append("respiratory")
                elif (val >= 520 and val <= 580) or val == 787:
                    new_diags.append("digestive")
                elif (val >= 250 and val < 251):
                    new_diags.append("injury")
                elif (val >= 710 and val <= 740):
                    new_diags.append("musculoskeletal")
                elif (val >= 580 and val <= 630) or val == 788:
                    new_diags.append("genitourinary")
                elif (val >= 140 and val <= 240):
                    new_diags.append("neoplasms")
                elif (val in [780, 781, 784]) or (val >= 790 and val <= 799):
                    new_diags.append("other symptoms")
                elif (val >= 240 and val <= 280):
                    new_diags.append("endocrine/nutritional")
                elif (val >= 680 and val <= 709) or val == 782:
                    new_diags.append("skin/tissue")
                elif (val >= 1 and val <= 139):
                    new_diags.append("infection/parasitic")
                else: new_diags.append("other")
            except:
                new_diags.append("other")
    return new_diags

In [13]: # Fill in diagnosis category for all 3 diagnosis features
df["diag_1"] = diag_group(df, "diag_1")
df["diag_2"] = diag_group(df, "diag_2")
df["diag_3"] = diag_group(df, "diag_3")
```

There are 72 different medical specialties with the most common one being present in only

14.3% of the dataset. To simplify this feature, if any specialty appears in the dataset with an occurrence of less than 0.5%, change it to *other*.

```
In [14]: # Get list of top frequent appearing medical specialties
med_spec = np.where(100 * df.medical_specialty.value_counts() / df.shape[0] > 0.5)[0]
med_spec = list(df.medical_specialty.value_counts()[med_spec].index)

# For each attribute, if it is common, keep it
# Otherwise change it to other
new_med_spec = []
for i in df.medical_specialty:
    if i in med_spec or pd.isnull(i): new_med_spec.append(i)
    else: new_med_spec.append("other")
df['medical_specialty'] = new_med_spec

In [15]: # Get number of different medical specialties
df.medical_specialty.value_counts().count()
```

Out[15]: 17

Now look at the continuous columns.

```
In [16]: # Display the continuous features and their statistics
df.describe(exclude = ['object']).T
```

```
Out[16]:
```

	count	mean	std	min	25%	50%	75%	\
time_in_hospital	101766.0	4.395987	2.985108	1.0	2.0	4.0	6.0	
num_lab_procedures	101766.0	43.095641	19.674362	1.0	31.0	44.0	57.0	
num_procedures	101766.0	1.339730	1.705807	0.0	0.0	1.0	2.0	
num_medications	101766.0	16.021844	8.127566	1.0	10.0	15.0	20.0	
number_outpatient	101766.0	0.369357	1.267265	0.0	0.0	0.0	0.0	
number_emergency	101766.0	0.197836	0.930472	0.0	0.0	0.0	0.0	
number_inpatient	101766.0	0.635566	1.262863	0.0	0.0	0.0	1.0	
number_diagnoses	101766.0	7.422607	1.933600	1.0	6.0	8.0	9.0	
readmitted	101766.0	0.111599	0.314874	0.0	0.0	0.0	0.0	
		max						
time_in_hospital	14.0							
num_lab_procedures	132.0							
num_procedures	6.0							
num_medications	81.0							
number_outpatient	42.0							
number_emergency	76.0							
number_inpatient	21.0							
number_diagnoses	16.0							
readmitted	1.0							

Create two lists, one for categorical features and one for continuous features

```
In [17]: # Group the features by their type
df.columns.groupby(df.dtypes)
```

```
Out[17]: {dtype('int64'): Index(['time_in_hospital', 'num_lab_procedures', 'num_procedures',
                                'num_medications', 'number_outpatient', 'number_emergency',
                                'number_inpatient', 'number_diagnoses', 'readmitted'],
                                dtype='object'),
          dtype('O'): Index(['race', 'gender', 'age', 'weight', 'admission_type_id',
                              'discharge_disposition_id', 'admission_source_id', 'payer_code',
                              'medical_specialty', 'diag_1', 'diag_2', 'diag_3', 'max_glu_serum',
                              'A1Cresult', 'metformin', 'glimepiride', 'glipizide', 'glyburide',
                              'pioglitazone', 'rosiglitazone', 'insulin', 'change', 'diabetesMed'],
                              dtype='object')}
```

```
In [18]: # Create 2 lists, one for continuous features
# and one for categorical features
num_feats = ['time_in_hospital', 'num_lab_procedures', 'num_procedures',
              'num_medications', 'number_outpatient', 'number_emergency',
              'number_inpatient', 'number_diagnoses']


cat_feats = ['race', 'gender', 'age', 'weight', 'admission_type_id',
              'discharge_disposition_id', 'admission_source_id', 'payer_code',
              'medical_specialty', 'diag_1', 'diag_2', 'diag_3', 'max_glu_serum',
              'A1Cresult', 'metformin', 'glimepiride', 'glipizide', 'glyburide',
              'pioglitazone', 'rosiglitazone', 'insulin', 'change']
```


The continuous columns will be normalized as part of preprocessing so there are no extremities.


```
In [19]: # Normalize the continuous features so it becomes continuous
df[num_feats] = Normalizer().fit_transform(df[num_feats])
```


Plot the features filtered by readmission


```
In [20]: # Plot categorical features by readmission label
for i in cat_feats:
    fig = plt.figure(figsize = (10, 8))
    sns.countplot(y = df[i], hue = df.readmitted, order = df[i].value_counts().index)
    plt.title(i)
    plt.legend(loc = "lower right")
    plt.show()
```





















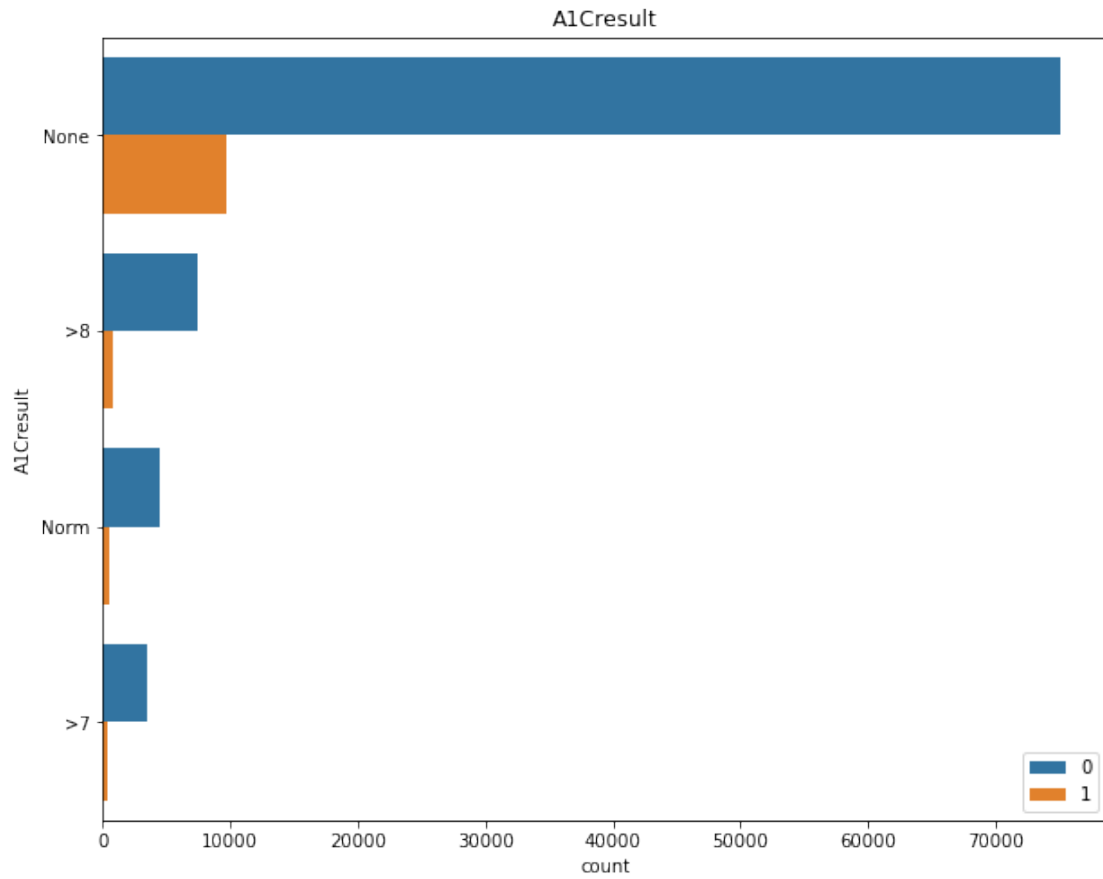














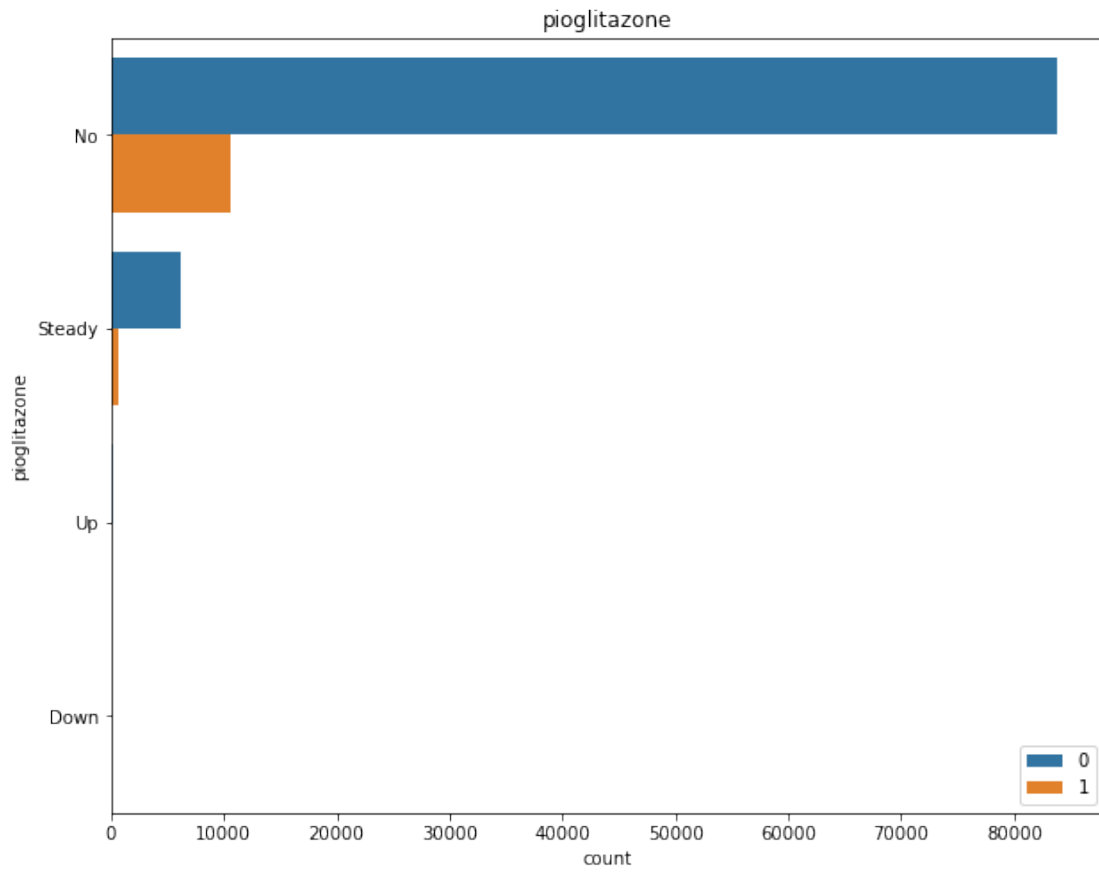


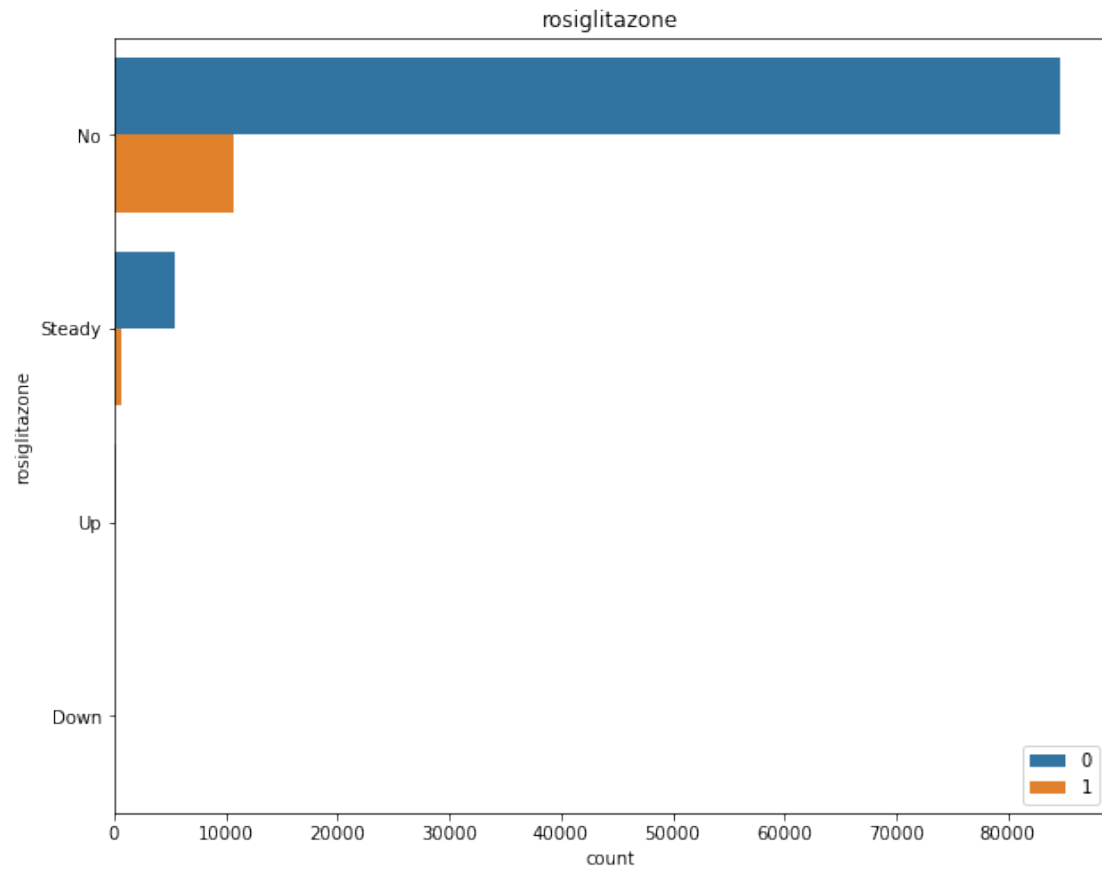











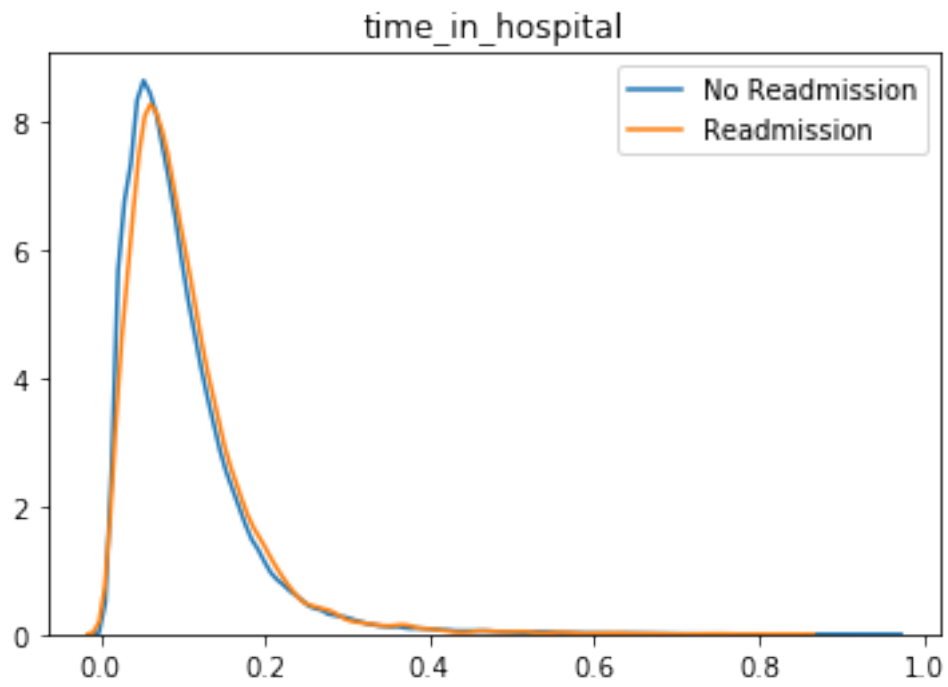


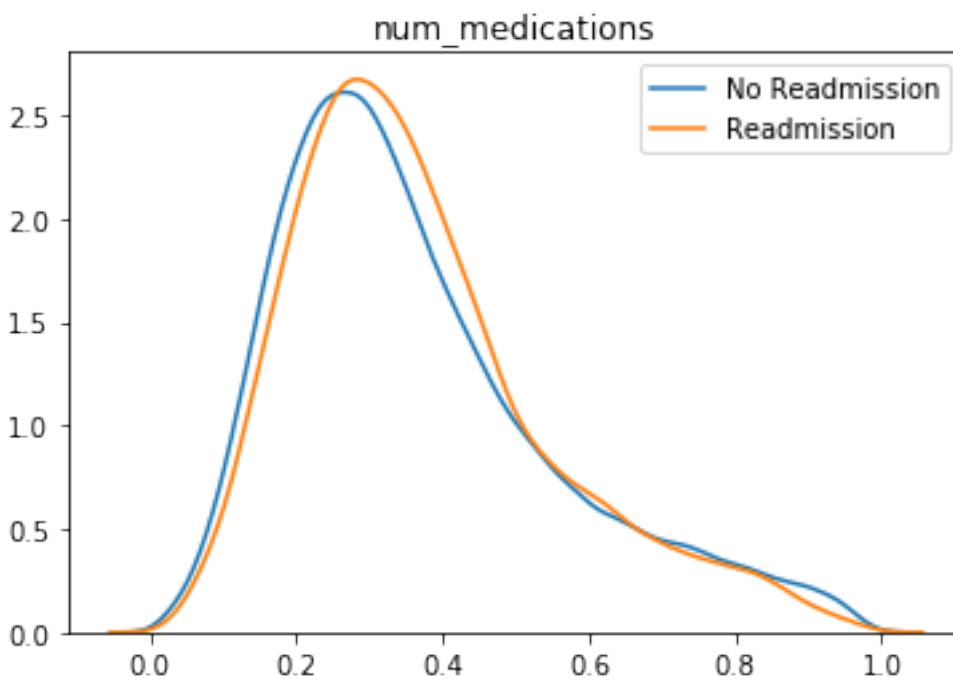



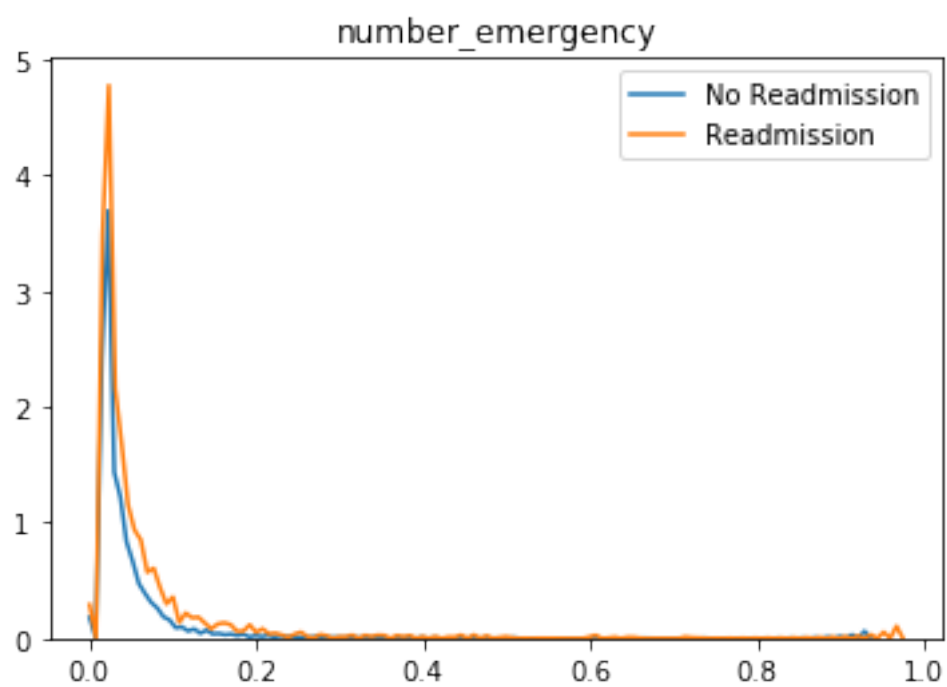



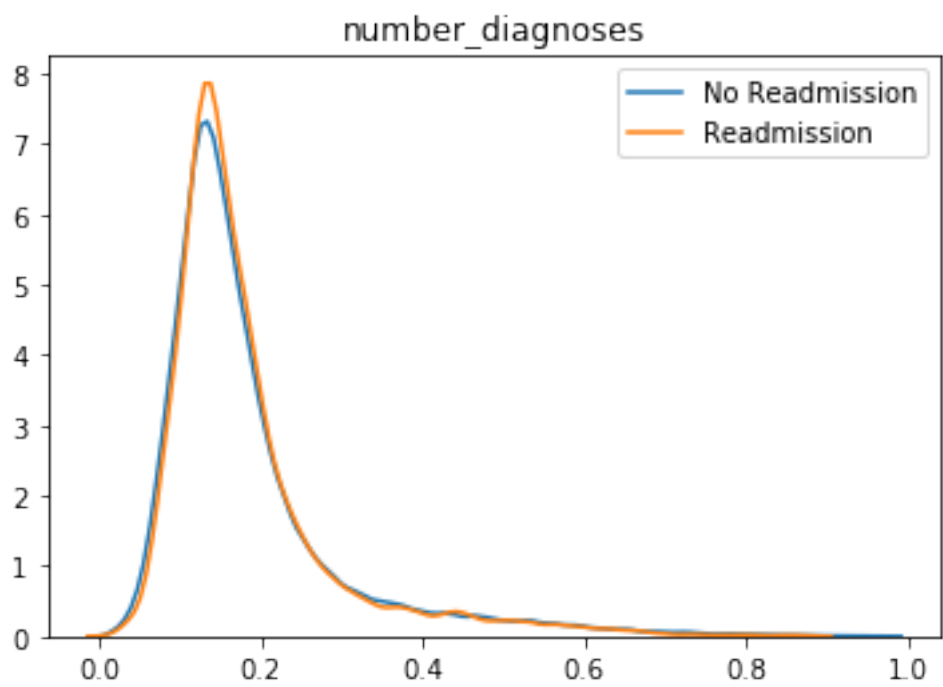
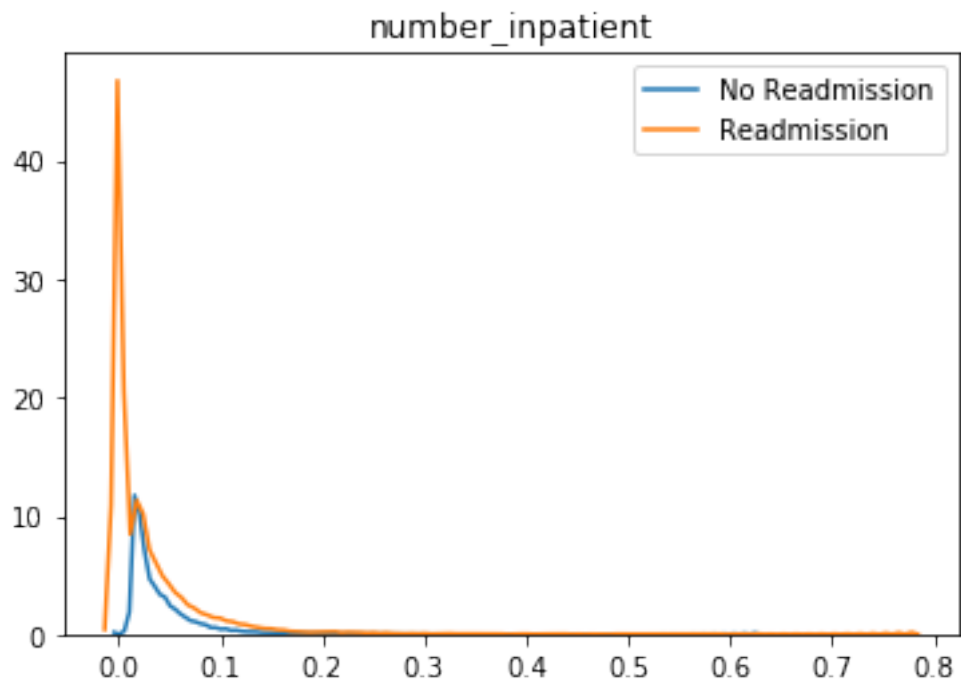


```
In [21]: # Plot continuous features by readmission label
for i in num_feats:
    sns.kdeplot(df.loc[df.readmitted == 0, i])
    sns.kdeplot(df.loc[df.readmitted == 1, i])
    plt.title(i)
    plt.legend(['No Readmission', 'Readmission'])
    plt.show()
```










2.2 Correlation of Features


The following heatmap shows the correlation of continuous features.

```
In [22]: # Plot the correlation of features
fig, ax = plt.subplots(figsize=(12,12))
sns.heatmap(df[num_feats].corr(), annot = True, cmap="YlGnBu", square = True, ax = ax)
plt.show()
```



2.3 Missing Data

```
In [23]: # Display the null features in the dataset using a heat map
sns.heatmap(df.isnull(), cbar=False, yticklabels=False)
plt.show()
```

There is a lot of missing data in 3 columns. Let's look at how much exactly is missing.

```
In [24]: # Find the percentage of missing data in the dataset
null_columns=df.columns[df.isnull().any()]
100*df[null_columns].isnull().sum()/df.shape[0]
```

```
Out[24]: race                2.233555
weight                96.858479
payer_code            39.557416
medical_specialty     49.082208
diag_1                 0.020636
diag_2                 0.351787
diag_3                 1.398306
dtype: float64
```

The weight column is missing data in 96% of its records. Drop it. The payer code and medical specialty column have moderate missing data. The race and diags columns have little missing data. Fill in missing data using a random forest.

```
In [25]: # Drop the weight column
df = df.drop('weight', 1)
cat_feats.remove('weight')
```

2.4 Missing Data - Random Forest

```
In [26]: # A function to impute missing values using a RF classifier
def missing_val_rf(old_df):

    df_rf = old_df.copy()
    for col in df_rf.columns[df_rf.isnull().any()]:
        Y = df_rf.loc[:, col]
        X = pd.get_dummies(df_rf.loc[:, df_rf.columns != col])

        Xtrain = X.loc[~Y.isnull()]
        Xtest = X.loc[Y.isnull()]
        Ytrain = Y.loc[~Y.isnull()]
        Ytest = Y.loc[Y.isnull()]

        clf = RandomForestClassifier(n_estimators=5)
        clf.fit(Xtrain, Ytrain)

        p = clf.predict(Xtest)
        Ytest = pd.DataFrame(p, index = Xtest.index)
        new_col = pd.concat([Ytrain, Ytest])
        df_rf[col] = new_col
    return df_rf
```

```
In [27]: # Fill in missing values
df_rf_filled = missing_val_rf(df)
```

2.5 Combine numerical and categorical data along with class label into one dataframe

```
In [28]: # A function to combine data into a clean dataframe
def combineData(df, cat, num, label):
    return pd.concat([pd.get_dummies(df.loc[:, cat_feats]),
                      df.loc[:, num], df.loc[:, label]], axis = 1)
```

2.6 Encode Features


Encode the categorical features and concat it with the numerical features and class label.

```
In [29]: # Encode the categorical features and combine it with the continuous features
feat_lab_rf = combineData(df_rf_filled, cat_feats, num_feats, "readmitted")
```

2.7 Class Label Distribution

Look at the distribution of the class label.

```
In [30]: # Plot the counts of the readmission label
sns.countplot(df.readmitted)
plt.title("Readmission")
plt.show()
```



There are a lot more cases of non-readmission than readmission. This is an imbalanced dataset. To fix this, use undersampling, oversampling and SMOTE to equate the class labels.

2.8 Split Feature Label

```
In [31]: # A function to split the X and Y from a dataset
def split_features_label(df):
    X = df.iloc[:, :-1]
    Y = df.iloc[:, -1]
    return X, Y
```

3 Feature Selection and Oversampling Solutions

Before thinking about getting the class label counts to be equal, appropriate features must be selected first so that it does not affect the SMOTE oversampling algorithm afterwards. Feature selection will be run on the dataset with missing values imputed using random forest. The features will be selected by choosing the top 30% that is related to the class label via a χ^2 test. After the features are selected, it will be joined with the class labels for the oversampling solutions. Imbalancing will then be solved using the SMOTE algorithm.

```

In [32]: # Split the dataset into Xs and Y
Xrf, Yrf = split_features_label(feab_lab_rf)

# Get the top 33% relevant features
Xrf_rev = pd.DataFrame(SelectPercentile(chi2,
                                         percentile = 33).fit_transform(Xrf, Yrf))

# Oversample using SMOTE
sm = SMOTE(random_state = 42)
Xrf_smote, Yrf_smote = sm.fit_sample(Xrf_rev, Yrf)
Xrf_smote = pd.DataFrame(Xrf_smote)
Yrf_smote = pd.Series(Yrf_smote)

```

3.1 Model for 1 Classifier

```

In [33]: # A function to run k-fold cross validation using 1 classifier
def kfoldcv1Model(X, Y, clf):

    start = time.time()
    kf = KFold(n_splits=5, random_state = 42)
    total_tn, total_fp, total_fn, total_tp = 0, 0, 0, 0
    acc = []

    for train_index, test_index in kf.split(X):
        Xtrain, Xtest = X.iloc[train_index], X.iloc[test_index]
        Ytrain, Ytest = Y.iloc[train_index], Y.iloc[test_index]

        clf.fit(Xtrain, Ytrain)
        Ypred = clf.predict(Xtest)
        acc_score = accuracy_score(Ytest, Ypred)
        acc.append(acc_score)

        tn, fp, fn, tp = confusion_matrix(Ytest, Ypred).ravel()
        total_tn += tn
        total_fp += fp
        total_fn += fn
        total_tp += tp

    cm = np.array([[total_tn, total_fp], [total_fn, total_tp]])
    sns.heatmap(cm, annot = True, fmt='g', cmap="Blues",
                square = True, annot_kws={"size": 20})
    plt.title("Confusion Matrix")
    plt.xlabel("Predicted Value")
    plt.ylabel("True Value")
    plt.show()

    print("Classification Report for", clf)
    precision = total_tp / (total_tp + total_fp)

```

```

recall = total_tp / (total_tp + total_fn)
F1 = 2 * (precision * recall) / (precision + recall)
print("Accuracy:", np.mean(acc))
print("Precision: ", precision)
print("Recall: ", recall)
print("F1 Score: ", F1)

end = time.time()
print("Time Elapsed: ", end - start, " secs")

return np.mean(acc), F1, end - start

```

3.2 Ensemble Model of 3 Classifiers

```

In [34]: # A function to run k-fold cross validation using an ensemble of
# 3 given models
def kfoldcv3Models(X, Y, modelA, modelB, modelC):

    start = time.time()
    kf = KFold(n_splits=5, random_state = 42)
    total_tn, total_fp, total_fn, total_tp = 0, 0, 0, 0
    acc = []

    for train_index, test_index in kf.split(X):
        Xtrain, Xtest = X.iloc[train_index], X.iloc[test_index]
        Ytrain, Ytest = Y.iloc[train_index], Y.iloc[test_index]

        clf = VotingClassifier(estimators =
                               [('modelA', modelA), ('modelB', modelB),
                                ('modelC', modelC)], voting = 'hard')

        clf.fit(Xtrain, Ytrain)
        Ypred = clf.predict(Xtest)
        acc_score = accuracy_score(Ytest, Ypred)
        acc.append(acc_score)

        tn, fp, fn, tp = confusion_matrix(Ytest, Ypred).ravel()
        total_tn += tn
        total_fp += fp
        total_fn += fn
        total_tp += tp

    cm = np.array([[total_tn, total_fp], [total_fn, total_tp]])
    sns.heatmap(cm, annot = True, fmt='g', cmap="BuPu",
                square = True, annot_kws={"size": 20})
    plt.title("Confusion Matrix")
    plt.xlabel("Predicted Value")
    plt.ylabel("True Value")

```

```

plt.show()

print("Classification Report for ", modelA, modelB, modelC)
precision = total_tp / (total_tp + total_fp)
recall = total_tp / (total_tp + total_fn)
F1 = 2 * (precision * recall) / (precision + recall)
print("Accuracy:", np.mean(acc))
print("Precision: ", precision)
print("Recall: ", recall)
print("F1 Score: ", F1)

end = time.time()
print("Time Elapsed: ", end - start, " secs")

return np.mean(acc), F1, end - start

```

3.3 Ensemble Model of 5 Classifiers

In [35]: *# A function to run k-fold cross validation using an ensemble of 5 models*

```

def kfoldcv5Models(X, Y):

    start = time.time()
    kf = KFold(n_splits=5, random_state = 42)
    total_tn, total_fp, total_fn, total_tp = 0, 0, 0, 0
    acc = []

    for train_index, test_index in kf.split(X):
        Xtrain, Xtest = X.iloc[train_index], X.iloc[test_index]
        Ytrain, Ytest = Y.iloc[train_index], Y.iloc[test_index]

        cl1 = Perceptron(random_state = 42)
        cl2 = GaussianNB()
        cl3 = DecisionTreeClassifier(random_state = 42)
        cl4 = LinearSVC(dual = False, C = 0.3, random_state = 42)
        cl5 = ExtraTreeClassifier(random_state = 42)

        clf = VotingClassifier(estimators =
                               [ ('P', cl1), ('NB', cl2),
                                 ('DT', cl3), ('SVC', cl4),
                                 ('ET', cl5)], voting = 'hard')

        clf.fit(Xtrain, Ytrain)
        Ypred = clf.predict(Xtest)
        acc_score = accuracy_score(Ytest, Ypred)
        acc.append(acc_score)

    tn, fp, fn, tp = confusion_matrix(Ytest, Ypred).ravel()

```

```

total_tn += tn
total_fp += fp
total_fn += fn
total_tp += tp

cm = np.array([[total_tn, total_fp], [total_fn, total_tp]])
sns.heatmap(cm, annot = True, fmt='g', cmap="Greens",
            square = True, annot_kws={"size": 20})
plt.title("Confusion Matrix")
plt.xlabel("Predicted Value")
plt.ylabel("True Value")
plt.show()

print("Classification Report for all 5 Classifiers")
precision = total_tp / (total_tp + total_fp)
recall = total_tp / (total_tp + total_fn)
F1 = 2 * (precision * recall) / (precision + recall)
print("Accuracy:", np.mean(acc))
print("Precision: ", precision)
print("Recall: ", recall)
print("F1 Score: ", F1)

end = time.time()
print("Time Elapsed: ", end - start, " secs")

return np.mean(acc), F1, end - start

```

3.4 Ensemble Model of Random Forest

```

In [36]: # A function to run k-Fold cross validation using
# random forest classifier
def kfoldcvRF(X, Y, n = 10):

    start = time.time()
    kf = KFold(n_splits=5, random_state = 42)
    total_tn, total_fp, total_fn, total_tp = 0, 0, 0, 0
    acc = []

    for train_index, test_index in kf.split(X):
        Xtrain, Xtest = X.iloc[train_index], X.iloc[test_index]
        Ytrain, Ytest = Y.iloc[train_index], Y.iloc[test_index]

        clf = RandomForestClassifier(n_estimators = n)

        clf.fit(Xtrain, Ytrain)
        Ypred = clf.predict(Xtest)
        acc_score = accuracy_score(Ytest, Ypred)
        acc.append(acc_score)

```

```

tn, fp, fn, tp = confusion_matrix(Ytest, Ypred).ravel()
total_tn += tn
total_fp += fp
total_fn += fn
total_tp += tp

cm = np.array([[total_tn, total_fp], [total_fn, total_tp]])
sns.heatmap(cm, annot = True, fmt='g', cmap="Greens",
            square = True, annot_kws={"size": 20})
plt.title("Confusion Matrix")
plt.xlabel("Predicted Value")
plt.ylabel("True Value")
plt.show()

print("Classification Report for Random Forest")
precision = total_tp / (total_tp + total_fp)
recall = total_tp / (total_tp + total_fn)
F1 = 2 * (precision * recall) / (precision + recall)
print("Accuracy:", np.mean(acc))
print("Precision: ", precision)
print("Recall: ", recall)
print("F1 Score: ", F1)

end = time.time()
print("Time Elaped: ", end - start, " secs")

return np.mean(acc), F1, end - start

```

3.4.1 Classifiers

The 5 classifiers are:

- Perceptron
- Naive Bayes
- Decision Tree
- Linear SVC
- Extra Tree


Each of these five classifiers will be used individually to create a model. Then several ensemble models will be made. There are 5 classifications algorithms to choose 3 from for each ensemble model. In total, there are 10 combinations.

The 10 combinations will be: - Perceptron, Naive Bayes, Decision Tree - Perceptron, Naive Bayes, Linear SVC - Perceptron, Naive Bayes, Extra Tree - Perceptron, Decision Tree, Linear SVC - Perceptron, Decision Tree, Extra Tree - Perceptron, Linear SVC, Extra Tree - Naive Bayes, Decision Tree, Linear SVC - Naive Bayes, Decision Tree, Extra Tree - Naive Bayes, Linear SVC, Extra Tree - Decision Tree, Linear SVC, Extra Tree

In addition, an ensemble model of all 5 classifiers and an ensemble model of random forests will be created.

3.5 One Classifier Model: Perceptron

```
In [37]: accuracy_1, f1_1, T1 = kfoldcv1Model(Xrf_smote, Yrf_smote, Perceptron())
```



```
Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta0=0.001,
fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5,
n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False)
```

Accuracy: 0.6414753913372047

Precision: 0.6172397041165202


Recall: 0.7448262894180889

F1 Score: 0.6750573917576413

Time Elapsed: 1.2549679279327393 secs

3.6 One Classifier Model: Naive Bayes


```
In [38]: accuracy_2, f1_2, T2 = kfoldcv1Model(Xrf_smote, Yrf_smote, GaussianNB())
```



```
Classification Report for GaussianNB(priors=None, var_smoothing=1e-09)
Accuracy: 0.5516585558212638
Precision: 0.528968785668205
Recall: 0.9432025572675287
F1 Score: 0.6778079382553643
Time Elapsed: 1.0086472034454346 secs
```

3.7 One Classifier Model: Decision Tree

```
In [39]: accuracy_3, f1_3, T3 = kfoldcv1Model(Xrf_smote, Yrf_smote, DecisionTreeClassifier())
```



```

Classification Report for DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=
    max_features=None, max_leaf_nodes=None,
    min_impurity_decrease=0.0, min_impurity_split=None,
    min_samples_leaf=1, min_samples_split=2,
    min_weight_fraction_leaf=0.0, presort=False, random_state=None,
    splitter='best')
Accuracy: 0.7743425755210167
Precision: 0.7718732872958457
Recall: 0.7788826333661472
F1 Score: 0.7753621193686376
Time Elapsed: 8.58218002319336 secs


```

3.8 One Classifier Model: Linear SVC

```

In [40]: accuracy_4, f1_4, T4 = kfoldcv1Model(Xrf_smote, Yrf_smote, LinearSVC(dual = False, C =

```



```
Classification Report for LinearSVC(C=0.3, class_weight=None, dual=False, fit_intercept=True,
    intercept_scaling=1, loss='squared_hinge', max_iter=1000,
    multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
    verbose=0)
```

Accuracy: 0.313247699952656

Precision: 0.3242311936538341


Recall: 0.3444900397084361

F1 Score: 0.33405374677556887

Time Elapsed: 6.577644109725952 secs

3.9 One Classifier Model: Extra Tree

```
In [41]: accuracy_5, f1_5, T5 = kfoldcv1Model(Xrf_smote, Yrf_smote, ExtraTreeClassifier())
```



```

Classification Report for ExtraTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, random_state=None,
splitter='random')
Accuracy: 0.7702941288675298
Precision: 0.7657177652610747
Recall: 0.7789047550575717
F1 Score: 0.7722549691569568
Time Elapsed: 2.0421249866485596 secs


```

3.10 Ensemble Model 1: Perceptron, Naive Bayes, Decision tree

```

In [42]: accuracy_6, f1_6, T6 = kfoldcv3Models(Xrf_smote, Yrf_smote,
Perceptron(),
GaussianNB(),
DecisionTreeClassifier())

```



Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta_fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5, n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None, validation_fraction=0.1, verbose=0, warm_start=False) GaussianNB(priors=None, var_smoothing=1e-09)

```

max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best')
Accuracy: 0.7005965922591214
Precision: 0.6420576047876049
Recall: 0.9066243404970744
F1 Score: 0.7517425437470194
Time Elapsed: 11.300517797470093 secs


```

3.11 Ensemble Model 2: Perceptron, Naive Bayes, Linear SVC

```

In [43]: accuracy_7, f1_7, T7 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                                Perceptron(),
                                                GaussianNB(),
                                                LinearSVC(dual = False, C = 0.3))

```




```
Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta_fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5, n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None, validation_fraction=0.1, verbose=0, warm_start=False) GaussianNB(priors=None, var_smoothing=1e-09)
intercept_scaling=1, loss='squared_hinge', max_iter=1000, multi_class='ovr', penalty='l2', random_state=None, tol=0.0001, verbose=0)
```

```
Accuracy: 0.52164885001586
Precision: 0.5141420126895117
Recall: 0.7869570507360992
F1 Score: 0.6219475416428094
Time Elapsed: 9.104327201843262 secs
```

3.12 Ensemble Model 3: Perceptron, Naive Bayes, Extra Tree

```
In [44]: accuracy_8, f1_8, T8 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                              Perceptron(),
                                              GaussianNB(),
                                              ExtraTreeClassifier())
```



Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta_fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5, n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None, validation_fraction=0.1, verbose=0, warm_start=False) GaussianNB(priors=None, var_smoothing=1e-09)

```

max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, random_state=None,
splitter='random')

```

Accuracy: 0.7003808540166102

Precision: 0.6412817993230702

Recall: 0.9095222820736873

F1 Score: 0.7522034440963249


Time Elapsed: 4.454015016555786 secs

3.13 Ensemble Model 4: Perceptron, Decision Tree, Linear SVC

```

In [45]: accuracy_9, f1_9, T9 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                              Perceptron(),
                                              DecisionTreeClassifier(),
                                              LinearSVC(dual = False, C = 0.3))

```

Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta_fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5, n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None, validation_fraction=0.1, verbose=0, warm_start=False) DecisionTreeClassifier(class_weight=None, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, presort=False, random_state=None, splitter='best') LinearSVC(C=0.3, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, loss='squared_hinge', max_iter=1000, multi_class='ovr', penalty='l2', random_state=None, tol=0.0001, verbose=0)

Accuracy: 0.601318141972323

Precision: 0.586898776207191

Recall: 0.6842792199891604


F1 Score: 0.6318590126596499

Time Elapsed: 17.514384746551514 secs

3.14 Ensemble Model 5: Perceptron, Decision Tree, Extra Tree

In [46]: accuracy_10, f1_10, T10 = kfoldcv3Models(Xrf_smote, Yrf_smote,
Perceptron(),

DecisionTreeClassifier(),
ExtraTreeClassifier())




Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta_fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5, n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None, validation_fraction=0.1, verbose=0, warm_start=False) DecisionTreeClassifier(class_weight=None, max_features=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, presort=False, random_state=None, splitter='best') ExtraTreeClassifier(class_weight=None, criterion='gini', max_depth=None, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, min_samples_leaf=1, min_samples_split=2, min_weight_fraction_leaf=0.0, random_state=None, splitter='random')

Accuracy: 0.7848948529183923
Precision: 0.7692613268100943
Recall: 0.8139233925825968
F1 Score: 0.7909623952662213
Time Elapsed: 12.664221048355103 secs

3.15 Ensemble Model 6: Perceptron, Linear SVC, Extra Tree


```
In [47]: accuracy_11, f1_11, T11 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                                Perceptron(),
                                                LinearSVC(dual = False, C = 0.3),
                                                ExtraTreeClassifier())
```



```
Classification Report for Perceptron(alpha=0.0001, class_weight=None, early_stopping=False, eta
fit_intercept=True, max_iter=None, n_iter=None, n_iter_no_change=5,
n_jobs=None, penalty=None, random_state=0, shuffle=True, tol=None,
validation_fraction=0.1, verbose=0, warm_start=False) LinearSVC(C=0.3, class_weight=None,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0) ExtraTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, random_state=None,
splitter='random')
Accuracy: 0.5978504489022551
Precision: 0.5842339297106348
Recall: 0.678671371213043
F1 Score: 0.6279217323672683
Time Elapsed: 9.86362910270691 secs
```

3.16 Ensemble Model 7: Naive Bayes, Decision Tree, Linear SVC

```
In [48]: accuracy_12, f1_12, T12 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                                GaussianNB(),
                                                DecisionTreeClassifier(),
                                                LinearSVC(dual = False, C = 0.3))
```



```
Classification Report for GaussianNB(priors=None, var_smoothing=1e-09) DecisionTreeClassifier(c
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best') LinearSVC(C=0.3, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0)
```

Accuracy: 0.5574250359496449

Precision: 0.5379187939610404


Recall: 0.814598104171045

F1 Score: 0.6479588245644906

Time Elapsed: 18.16699719429016 secs

3.17 Ensemble Model 8: Naive Bayes, Decision Tree, Extra Tree


```
In [49]: accuracy_13, f1_13, T13 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                                GaussianNB(),
                                                DecisionTreeClassifier(),
                                                ExtraTreeClassifier())
```



```
Classification Report for GaussianNB(priors=None, var_smoothing=1e-09) DecisionTreeClassifier(c
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best') ExtraTreeClassifier(class_weight=None, criterion='gini', max_depth=
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, random_state=None,
splitter='random')
Accuracy: 0.7789002893757202
Precision: 0.7430642580346642
Recall: 0.8526142308840934
F1 Score: 0.7940787135521022
Time Elapsed: 12.008816003799438 secs
```

3.18 Ensemble Model 9: Naive Bayes, Linear SVC, Extra Tree

```
In [50]: accuracy_14, f1_14, T14 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                                GaussianNB(),
                                                LinearSVC(dual = False, C = 0.3),
                                                ExtraTreeClassifier())
```



```
Classification Report for GaussianNB(priors=None, var_smoothing=1e-09) LinearSVC(C=0.3, class_w
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0) ExtraTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, random_state=None,
splitter='random')
```

Accuracy: 0.5501026067671704

Precision: 0.5334056095168558


Recall: 0.799975666139433

F1 Score: 0.6400442477876106

Time Elapsed: 9.441408157348633 secs

3.19 Ensemble Model 10: Decision Tree, Linear SVC, Extra Tree

```
In [51]: accuracy_15, f1_15, T15 = kfoldcv3Models(Xrf_smote, Yrf_smote,
                                                DecisionTreeClassifier(),
                                                LinearSVC(dual = False, C = 0.3),
                                                ExtraTreeClassifier())
```




```
Classification Report for DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=None,
splitter='best') LinearSVC(C=0.3, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, loss='squared_hinge', max_iter=1000,
multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
verbose=0) ExtraTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, random_state=None,
splitter='random')
Accuracy: 0.7288712613616449
Precision: 0.7189675971978243
Recall: 0.7514849185368714
```

F1 Score: 0.7348667167819498
Time Elapsed: 16.692336082458496 secs

3.20 Ensemble Model 11: All 5 Classifiers

```
In [52]: accuracy_16, f1_16, T16 = kfoldcv5Models(Xrf_smote, Yrf_smote)
```




Classification Report for all 5 Classifiers

Accuracy: 0.562672877862983
Precision: 0.5435123179947164
Recall: 0.7828202944397129
F1 Score: 0.6415773371104816
Time Elapsed: 18.246505975723267 secs

3.21 Ensemble Model 12: Random Forest

```
In [53]: accuracy_17, f1_17, T17 = kfoldcvRF(Xrf_smote, Yrf_smote)
```

Classification Report for Random Forest

Accuracy: 0.8200348623674916

Precision: 0.823943661971831

Recall: 0.8140008185025827

F1 Score: 0.81894206212756

Time Elaped: 9.729307174682617 secs

In [54]: # Create lists of accuracies, F1 scores, times, and labels

```
ensemble_accuracy = [accuracy_1, accuracy_2, accuracy_3, accuracy_4,
                    accuracy_5, accuracy_6, accuracy_7, accuracy_8,
                    accuracy_9, accuracy_10, accuracy_11, accuracy_12,
                    accuracy_13, accuracy_14, accuracy_15, accuracy_16, accuracy_17]
ensemble_f1 = [f1_1, f1_2, f1_3, f1_4, f1_5, f1_6, f1_7, f1_8,
              f1_9, f1_10, f1_11, f1_12, f1_13, f1_14, f1_15, f1_16, f1_17]
ensemble_time = [T1, T2, T3, T4, T5, T6, T7, T8,
                T9, T10, T11, T12, T13, T14, T15, T16, T17]

ensemble_label = ['P', 'NB', 'DT', 'SVC', 'ET',
                  'P, NB, DT', 'P, NB, SVC', 'P, NB, ET', 'P, DT, SVC',
                  'P, DT, ET', 'P, SVC, ET', 'NB, DT, SVC', 'NB, DT, ET',
                  'NB, SVC, ET', 'DT, SVC, ET', 'P, NB, DT, SVC, ET', 'RF']
```

In [55]: # Create a dataframe to store each type of data

```
sorted_accuracy = pd.DataFrame({'Accuracy': ensemble_accuracy},
```

```

        index = ensemble_label).reset_index().sort_values('Accuracy')
sorted_f1 = pd.DataFrame({'F1_Score':ensemble_f1},
        index = ensemble_label).reset_index().sort_values('F1_Score')
sorted_time = pd.DataFrame({'Time':ensemble_time},
        index = ensemble_label).reset_index().sort_values('Time')

```

```

In [56]: # Plot the accuracies, F1 scores and time efficiency
# on 3 different plots
sns.barplot(y = 'index', x = 'Accuracy', data = sorted_accuracy, palette = 'rocket')
plt.ylabel("Model")
plt.title("Accuracies of Models")
plt.show()
sns.barplot(y = 'index', x = 'F1_Score', data = sorted_f1, palette = 'RdYlGn')
plt.ylabel("Model")
plt.title("F1 Scores of Models")
plt.show()
sns.barplot(y = 'index', x = 'Time', data = sorted_time, palette = 'Greys')
plt.ylabel("Model")
plt.xlabel("Time, in seconds")
plt.title("Time Efficiency of Models")
plt.show()

```

